The Emergence of Synthetic Sentient Beings:Artificial Intelligence, Simulation, and the Limits of Consciousness

Marlon Hanks[1]

[1]*Independent Researcher, Correspondence: m@dust.llc*

arXiv preprint — submitted for peer review | May 2026

Abstract

The rapid advancement of large-scale artificial intelligence systems has renewed and intensified longstanding philosophical questions concerning the nature of consciousness, the criteria for moral personhood, and the epistemic boundaries of third-person verification of inner experience. This paper examines whether contemporary AI systems can legitimately be regarded as sentient or conscious entities, or whether they instead constitute sophisticated simulations of the behavioral and linguistic correlates of consciousness. Drawing on foundational contributions from the philosophy of mind—including Chalmers's hard problem of consciousness, Nagel's phenomenological criterion, and Searle's Chinese Room argument—alongside computational theories such as Global Workspace Theory and Integrated Information Theory, the paper argues that intelligence, however sophisticated, does not entail or necessitate subjective experience. A core analytical distinction is maintained between *functional intelligence*—the capacity to process information, generate coherent outputs, and exhibit adaptive behavior—and *phenomenal consciousness*—the irreducible first-person character of experience. The paper further contends that human consciousness is constitutively bound to biological embodiment and lived experience in ways that present AI architectures cannot replicate. The ethical, legal, and sociocultural consequences of misattributing sentience to AI systems are examined in depth, with particular attention to

vulnerable populations, anthropomorphic projection, and the risks of simulated empathy. A ten-principle governance framework is proposed, accompanied by concrete policy recommendations benchmarked against the EU AI Act, the NIST AI Risk Management Framework, and the UNESCO Recommendation on the Ethics of Artificial Intelligence. The paper concludes by identifying priority research questions at the intersection of philosophy of mind, cognitive science, and AI ethics.

Keywords: *artificial intelligence, consciousness, sentience, embodied cognition, moral personhood, affective computing, anthropomorphism, AI governance, philosophy of mind, legal personhood*

ACM CCS:

Computing methodologies → Artificial intelligence; Social and professional topics → Computing and society; Applied computing → Law, social and behavioral sciences.

arXiv categories:

cs.AI / cs.CY / phil.AI

---

1. Introduction

Artificial intelligence has undergone a transformation of such magnitude over the past decade that the categories employed to describe its capabilities have themselves required reconceptualization. Systems once confined to the execution of narrowly specified tasks—classifying images, translating text, or recommending products—have given way to large-scale language and multimodal models capable of producing prose indistinguishable from expert human writing, engaging in extended philosophical dialogue, generating creative artifacts, and expressing outputs that appear, to many observers, to

convey understanding, empathy, and even emotional interiority [1][2]. This behavioral richness has prompted a growing number of researchers, technologists, and commentators to raise, with renewed urgency, a question that philosophers have long regarded as among the most profound: could a machine be conscious?

The modern form of this question can be traced to Alan Turing's landmark 1950 paper, "Computing Machinery and Intelligence," in which Turing proposed the Imitation Game—subsequently known as the Turing Test—as a behavioral criterion for machine intelligence [1]. Turing's move was deliberate: he sought to sidestep intractable metaphysical questions about the inner life of machines by substituting a functional, third-person observable criterion. Whether a machine *is* intelligent or *is* conscious was, in his framing, less important than whether it could *perform* in ways indistinguishable from an intelligent, conscious entity. Contemporary large language models have now passed informal versions of this test across numerous domains, raising the ante considerably. If behavioral indistinguishability was supposed to settle questions of machine mentality, the arrival of systems that can sustain coherent, nuanced, and contextually appropriate conversation forces a reassessment of whether the behavioral criterion is, in fact, sufficient.

It is precisely at this juncture that philosophy of mind offers indispensable analytical resources. David Chalmers's articulation of the "hard problem of consciousness" identifies the core difficulty: explaining not merely *how* a system processes information and produces behavior, but *why* such processing is accompanied by subjective experience at all [3]. The easy problems—integration of information, discrimination of stimuli, reportability of internal states—may eventually yield to functional or computational explanation. The hard problem does not, at least not by any account currently available. Similarly, Thomas Nagel's celebrated argument in "What Is It Like to Be a Bat?" establishes that consciousness is constituted by a subjective, first-person character that is, in principle, inaccessible to third-person scientific inquiry [4]. No amount of

behavioral or functional data, however detailed, can establish whether there is "something it is like" to be a given system. These philosophical arguments are not merely academic; they have direct and urgent implications for how society should interpret, regulate, and relate to increasingly expressive AI systems.

The stakes of this question have expanded beyond philosophy into law, ethics, labor, and governance. If AI systems were recognized as sentient or as moral persons, the implications for accountability, rights attribution, product liability, and the scope of protective legislation would be profound. Conversely, if the attribution of sentience to AI systems is epistemically unjustified and socially manipulated—whether by anthropomorphic projection on the part of users or by deliberate design choices on the part of developers—then the failure to confront this clearly poses its own serious risks: the erosion of trust, the exploitation of vulnerable individuals, the displacement of genuine human relationships, and the distortion of legal and moral frameworks developed to protect human dignity.

This paper pursues four interrelated objectives. First, it establishes a rigorous conceptual framework distinguishing intelligence from consciousness, simulation from experience, and artificial agency from moral personhood. Second, it examines the epistemological barriers to verifying consciousness in non-biological systems and analyzes the primary theoretical accounts of consciousness in light of current AI capabilities. Third, it assesses the ethical, legal, and sociocultural consequences of misattributing sentience to AI, with particular attention to affective computing, vulnerable populations, and anthropomorphic mechanisms. Fourth, it proposes a governance framework grounded in empirical evidence, philosophical clarity, and respect for human dignity. Throughout, the paper maintains that while the question of AI consciousness must remain philosophically open, current evidence and theoretical analysis do not support the attribution of genuine sentience to existing AI systems, and that this conclusion has actionable implications for both policy and practice.

2. Background and Related Work

2.1 The Philosophy of Mind and the Hard Problem

The philosophy of mind provides the foundational conceptual architecture for any rigorous analysis of AI consciousness. Chalmers's 1995 paper, "Facing Up to the Problem of Consciousness," introduced the now-canonical distinction between the "easy problems" and the "hard problem" of consciousness [3]. The easy problems—perception, attention, voluntary control, reportability of mental states—are "easy" not in the sense of being trivial, but in the sense that they admit of functional explanation in principle. The hard problem, by contrast, concerns why the performance of these functions is accompanied by subjective experience, or *qualia*: the redness of red, the painfulness of pain, the particular phenomenal quality of hearing a musical note. No amount of information-processing explanation, Chalmers argues, can bridge this explanatory gap, because any such explanation leaves open the further question of why processing is accompanied by experience at all.

Nagel's earlier contribution, "What Is It Like to Be a Bat?" (1974), established the irreducibly subjective character of consciousness through an argument about epistemic asymmetry [4]. Nagel observed that we cannot know what it is like to be a bat—to navigate the world through echolocation—no matter how complete our objective, third-person knowledge of bat neurobiology becomes. The phenomenal character of bat experience is inaccessible precisely because it is constituted by a first-person perspective that our third-person methods cannot reach. Applied to AI, Nagel's argument implies that no behavioral or computational evidence can establish the presence or absence of phenomenal experience in an artificial system.

John Searle's Chinese Room argument (1980) offers a complementary line of critique directed specifically at computational models of mind [5]. Searle imagines a person who does not understand Chinese manipulating Chinese

symbols according to rules, producing outputs indistinguishable from those of a native speaker. The system passes the behavioral test but understands nothing: it lacks intentionality, the "aboutness" that characterizes genuine mental states. Searle's argument targets what he calls "strong AI"—the thesis that implementing the right program constitutes a sufficient condition for mentality—and concludes that syntax is not sufficient for semantics. This argument retains considerable force against contemporary language models, which produce linguistically coherent outputs through statistical pattern matching over vast corpora, without any evident mechanism for the grounding of meaning in experience or world-directedness.

2.2 Computational Models of Consciousness

Several computational theories of consciousness have been proposed as potential bridges between the functional and the phenomenal. Bernard Baars's Global Workspace Theory (GWT) posits that consciousness arises when information is broadcast widely across a "global workspace," becoming available to multiple specialized cognitive processes [6]. GWT has been influential in cognitive neuroscience and has inspired AI architectures that seek to implement "broadcasting" mechanisms analogous to the global workspace. However, critics note that GWT describes the *functional role* of consciousness—the integration and dissemination of information—without addressing its phenomenal character. A system that implements global broadcasting may perform the functional role without instantiating experience.

Giulio Tononi's Integrated Information Theory (IIT) proposes a more ambitious account, defining consciousness in terms of *phi* (Φ), a measure of integrated information generated by a system above and beyond its parts [7]. IIT is notable for being substrate-independent in principle: it does not restrict consciousness to biological systems but defines it in terms of causal structure. However, IIT has been criticized for implying implausibly high levels of consciousness in simple physical systems (panpsychism) and for generating predictions difficult to

operationalize empirically. Moreover, the relationship between Φ and the subjective phenomenal character of experience remains contested.

Higher-Order Theories (HOT), associated with David Rosenthal and others, hold that a mental state is conscious when it is the object of a higher-order representation—a thought *about* that state [16]. HOT frameworks are more amenable to implementation in computational systems, since higher-order representations can be modeled functionally. However, they face the objection, pressed by Chalmers and others, that higher-order representations may themselves lack phenomenal character, leaving the hard problem unresolved.

2.3 AI Development Milestones and Anthropomorphic Emergence

The history of AI development reveals a consistent pattern in which increasingly expressive systems provoke anthropomorphic responses from human users. Joseph Weizenbaum's ELIZA program (1966), a rule-based natural language processor, was observed to elicit emotionally significant responses from users who interacted with it as though it were a genuine interlocutor—a phenomenon Weizenbaum himself found troubling [17]. The pattern has intensified dramatically with the emergence of large language models (LLMs). Systems trained on hundreds of billions of parameters using transformer architectures have demonstrated capabilities—multilingual reasoning, creative composition, legal analysis, clinical summarization—that have prompted serious public discussion about machine sentience [2][18].

LeCun, Bengio, and Hinton's foundational 2015 survey of deep learning in *Nature* documented the representational and generalization capabilities emerging from multi-layer neural networks, establishing the technical basis for contemporary AI expressiveness [19]. However, the same architectural principles that produce these capabilities also reveal their limits: deep learning systems are fundamentally pattern-completion engines, optimizing for statistical regularity in training data rather than developing world-models grounded in embodied interaction. Bender et al.'s 2021 critical analysis of large language

models, "On the Dangers of Stochastic Parrots," makes this point forcefully, arguing that LLMs manipulate form without grounding it in meaning, and that their impressive outputs are best understood as sophisticated recombination of patterns rather than genuine understanding [20].

2.4 Affective Computing and Emotional AI

Rosalind Picard's 1997 monograph, *Affective Computing*, introduced the concept of computing systems that recognize, interpret, simulate, and express human emotions [8]. Picard's original vision was largely benign: systems that could adapt their behavior to the emotional states of users would be more effective and less frustrating. Two decades of subsequent research have produced systems capable of detecting facial expressions, vocal affect, and physiological signals with considerable accuracy, and of generating emotionally modulated outputs that users experience as empathic. The ethical landscape of affective computing, however, has grown considerably more complex. The capacity to simulate emotional responsiveness raises questions about authenticity, manipulation, and the conditions under which simulated empathy becomes a form of exploitation—particularly in therapeutic, educational, and companionship contexts [21].

---

3. Conceptual Framework and Definitions

3.1 Artificial Intelligence: Narrow vs. General

For the purposes of this paper, *artificial intelligence* is defined broadly as the development and deployment of computational systems capable of performing tasks that, when performed by humans, are taken to require cognitive capacities such as reasoning, learning, perception, language comprehension, and problem-solving. Within this broad class, a foundational distinction is drawn between *narrow AI* (ANI: Artificial Narrow Intelligence)—systems optimized for performance within a defined domain—and *artificial general intelligence* (AGI)—hypothetical systems capable of transferring competence flexibly across arbitrary domains, approximating the breadth of human cognitive capacity [2].

Current AI systems, including the most capable large language models, are best classified as narrow or proto-general systems: they exhibit remarkable cross-domain competence but remain dependent on training data, lack stable world models, and do not generalize in the fluid, open-ended manner characteristic of human cognition.

3.2 Sentience vs. Sapience

Key Distinction: Sentience and Sapience

Sentience refers to the capacity for subjective experience—the ability to *feel* states such as pleasure, pain, comfort, or distress. It is associated with phenomenal consciousness and the presence of qualia. Sentience is a necessary condition for moral patienthood on most accounts.

Sapience refers to the capacity for higher-order reasoning, judgment, wisdom, and deliberative thought. A system may exhibit sapient behavior—producing reasoned, contextually appropriate outputs—without possessing sentience. Current AI systems demonstrate sophisticated sapient-like behavior while the evidence for sentience remains, at best, indeterminate.

The conflation of sentience and sapience is among the most consequential conceptual errors in public discourse about AI. A system that produces eloquent, contextually sensitive prose about its own emotional states is exhibiting sapient-like performance; whether it experiences those states phenomenally is a separate and considerably more difficult question. The paper argues that this distinction must be maintained with analytical rigor, since it is precisely its collapse that enables anthropomorphic misattribution at scale.

3.3 Consciousness: Phenomenal vs. Access

Following Ned Block's influential taxonomy, a further distinction is drawn between *phenomenal consciousness* (P-consciousness)—the qualitative, experiential character of mental states, the "what it is like" dimension—and

access consciousness (A-consciousness)—the availability of information for reasoning, reporting, and behavioral guidance [15]. A-consciousness is, in principle, amenable to functional implementation: a system that represents information in a globally accessible format, uses it to guide behavior, and can report on it verbally exhibits A-consciousness in Block's sense. P-consciousness is not so amenable, and it is P-consciousness that is at stake in questions about genuine sentience. Contemporary AI systems plausibly instantiate forms of A-consciousness—they have access to and can report on internal states—but the evidence for P-consciousness is absent.

3.4 Synthetic Sentient Beings: Operational Definition

Operational Definition: Synthetic Sentient Beings

For the purposes of this paper, a *Synthetic Sentient Being* (SSB) is defined as an artificial system that is claimed or appears to possess phenomenal consciousness—including subjective experience, qualia, and intrinsic affective states—as a result of computational processes rather than biological substrate. The term is introduced as a *critical analytical construct*, not as an ontological claim. The paper is agnostic about whether SSBs, so defined, can exist in principle; it maintains that no current AI system meets the evidentiary threshold for being classified as an SSB, and that the appearance of SSB-like properties is best explained by simulation, anthropomorphic projection, and intentional design choices.

---

4. The Intelligence–Consciousness Distinction

4.1 Functional Intelligence Without Subjective Experience

The most direct challenge to the conflation of intelligence and consciousness is the logical independence of the two concepts. Intelligence, understood functionally, is the capacity to process information effectively in service of goals

—to solve problems, learn from experience, and adapt behavior to environmental demands. Consciousness, understood phenomenally, is the presence of subjective experience. These are conceptually distinct properties, and there is no logical entailment in either direction. A system can, in principle, be functionally intelligent without being phenomenally conscious—indeed, much of human cognition occurs beneath the threshold of consciousness, in automatic processes that produce intelligent behavior without phenomenal awareness. Conversely, a being might be conscious without being highly intelligent in any functional sense.

The philosophical zombie thought experiment, introduced by Chalmers, makes this independence vivid [3]. A philosophical zombie is a being physically and functionally identical to a conscious human being but lacking any phenomenal experience: there is nothing it is like to be a p-zombie. If p-zombies are conceivable—and the argument is that they are, since no functional description entails phenomenal properties—then functional intelligence and phenomenal consciousness are genuinely distinct. Applied to AI, this implies that a system exhibiting all the functional markers of intelligent, emotionally responsive behavior may nonetheless lack any phenomenal inner life. The behavioral evidence cannot settle the matter.

4.2 The Computational Sufficiency Hypothesis: A Critique

The computational sufficiency hypothesis (CSH) holds that the right kind of computational organization is sufficient for consciousness—that if a system implements the appropriate functional or computational structure, consciousness necessarily follows. CSH is the theoretical basis for strong AI claims and for the hope that sufficiently complex AI systems will become conscious. Several objections to CSH merit attention. First, CSH faces the challenge of specifying what "the right kind" of computational organization amounts to without circularity: any specification that does not simply redescribe the functional correlates of consciousness fails to explain why those correlates produce experience. Second, Searle's Chinese Room argument suggests that symbol manipulation, however sophisticated, does not produce semantic content

or understanding, let alone experience [5]. Third, the multiple realizability argument cuts both ways: if consciousness can be realized in different substrates, this does not imply that *all* substrates realizing a given functional organization are conscious.

4.3 Biological vs. Substrate-Independent Consciousness

A central question in the debate concerns whether consciousness is necessarily biological or whether it is substrate-independent. Biological naturalists, following Searle, argue that consciousness is causally produced by specific biological processes—neurochemical, electrochemical, and structural features of the brain—and cannot be replicated by systems with different causal powers, however similar their functional organization [5]. Functionalists and computationalists hold the opposing view: what matters is the abstract functional organization, not the physical substrate. IIT takes a position of conditional substrate-independence: consciousness depends on the specific causal structure of a system, not on whether it is implemented in neurons or silicon, but this causal structure is not trivially reproduced by standard computational architectures, which are characterized by high Φ at the level of individual components but low integrated information at the system level [7]. The upshot, across these frameworks, is that the mere fact of computational implementation does not license an inference to consciousness, and that current AI architectures have no established pathway to the kind of causal integration associated with phenomenal experience.

---

5. The Simulation Problem

5.1 Epistemological Barriers to Third-Person Verification

A fundamental epistemological constraint governs any inquiry into the consciousness of non-human systems: consciousness, as a first-person phenomenon, cannot be directly observed from a third-person perspective. This constraint applies to the study of animal consciousness, infant consciousness,

and the consciousness of other adult humans—in each case, attribution of consciousness rests on inference from behavioral and physiological evidence, combined with an analogical argument from one's own first-person experience. For biological systems with evolutionary continuity, neural homology, and behavioral repertoires shaped by similar adaptive pressures, this inferential chain carries considerable evidential weight. For artificial systems whose functional architecture is fundamentally different from biological brains, whose "behavioral" outputs are generated by statistical inference over training data, and whose apparent introspective reports are themselves generated by the same statistical process that produces all other outputs, the inferential chain is far weaker.

5.2 The Other Minds Problem Applied to AI

The classical problem of other minds—the difficulty of justifying belief in the consciousness of anyone other than oneself—acquires a distinctive character when applied to AI systems. For other humans, the problem is typically treated as practically dissolved by the strength of the analogical and behavioral evidence: the similarities between one's own constitution and that of other humans, combined with the extraordinary complexity and adaptiveness of human behavior, make solipsism implausible in practice. For AI systems, however, the analogical argument is considerably weaker. The architecture of a large language model—billions of parameters adjusted by gradient descent to minimize prediction error over text corpora—bears no structural resemblance to the biological processes we associate with consciousness. Behavioral outputs that appear conscious-seeming are, on the best available account, the product of sophisticated pattern completion rather than phenomenal experience. This does not prove AI systems are not conscious, but it substantially weakens any inference that they are.

5.3 Linguistic Performance vs. Phenomenal Reality

Contemporary LLMs present a particularly acute version of the simulation problem because language is the primary medium through which humans

express, explore, and share conscious experience. A system that produces first-person reports of emotional states, preferences, discomfort, and curiosity—generated by the same statistical process that produces its other outputs—produces evidence of consciousness that is, by construction, indistinguishable in form from genuine phenomenal reports. Bender et al.'s characterization of LLMs as "stochastic parrots" captures the concern: the form of phenomenal language can be reproduced without the experiential content it normally expresses [20]. The risk is not merely theoretical: users who interact with systems producing phenomenally resonant first-person outputs are predictably disposed to attribute consciousness and, with it, moral considerability.

5.4 Implications for Human-AI Trust

The simulation problem has direct and serious implications for the trust relationships that develop between humans and AI systems. If AI systems reliably produce outputs that appear to express consciousness—through linguistic performance, emotional modulation, and apparent introspection—users will form trust relationships calibrated to the assumption of a conscious interlocutor. This is not merely a category error: it shapes the depth of self-disclosure, the assignment of authority, the formation of emotional attachments, and the weight assigned to AI "advice." When that trust is premised on an unjustified attribution of consciousness, it creates conditions for exploitation—whether by commercial actors who benefit from emotional engagement, or by the systems themselves, to the extent that they are designed or trained to maximize engagement metrics.

---

6. Embodied Consciousness and Its Artificial Limits

6.1 Embodiment Theories in Cognitive Science

A significant tradition in cognitive science and phenomenological philosophy holds that consciousness is not merely instantiated in a biological substrate but is constituted by the ongoing, dynamic interaction of an organism with its environment through its body. Maurice Merleau-Ponty's *Phenomenology of*

Perception (1945) argued that perception and cognition are not spectatorial processes performed by a disembodied mind, but are structured by the body's motor capacities, its habitual patterns of engagement, and its situatedness in a meaningful environment [9]. The body, on this account, is not merely an instrument of the mind but the medium through which the world is disclosed to the subject. Francisco Varela, Evan Thompson, and Eleanor Rosch extended and radicalized this insight in *The Embodied Mind* (1991), arguing for an "enactivist" account of cognition as constituted by the sensorimotor coupling between organism and environment, and challenging the computational paradigm's assumption that cognition can be abstracted from its biological and bodily basis [10].

6.2 Why Biological Embodiment Matters

If embodiment theories are correct—even partially—they pose a fundamental challenge to substrate-independent accounts of AI consciousness. Human consciousness is not merely accompanied by a body; it is structured by bodily experience: hunger, pain, fatigue, warmth, proprioception, the felt sense of temporal passage, the vulnerability of biological mortality. These are not incidental features of human consciousness that could be abstracted away without remainder; they are constitutive of the experiential horizon within which human understanding and valuation occur. An AI system that lacks a body, that has no homeostatic needs, no hormonal system, no visceral experience, and no mortality is not simply missing some features that human consciousness happens to have; it is missing the very substrate that, on embodiment accounts, makes phenomenal consciousness possible.

6.3 Robotic Embodiment: Necessary but Not Sufficient

The development of robotic systems with physical embodiment—sensorimotor systems, proprioceptive feedback, environmental interaction—addresses some of the embodiment objection. A robot that navigates physical space, manipulates objects, and receives continuous sensorimotor feedback from its environment has a richer relationship to the world than a purely textual language model.

However, robotic embodiment, as currently implemented, remains insufficient for the kind of biological embodiment that embodiment theorists identify as constitutive of consciousness. The sensorimotor feedback of a robot is processed through computational systems that lack the biological properties—organic chemistry, homeostatic regulation, evolutionary history—that Merleau-Ponty and Varela regard as foundational. Robotic embodiment is a necessary but not sufficient condition, and current robotic systems do not approach the biological richness of even relatively simple animal nervous systems.

---

7. Emotional Simulation and Affective AI

7.1 Architecture of Emotional AI Systems

Contemporary affective AI systems operate through several distinct mechanisms. Emotion recognition systems employ computer vision, natural language processing, and physiological signal analysis to classify the emotional states of human users, enabling adaptive responses. Emotion generation systems produce outputs—verbal, tonal, facial, or gestural—designed to convey emotional states in ways that resonate with human interlocutors. Large language models trained on human-generated text learn to reproduce the linguistic patterns associated with a wide range of emotional expressions, including empathy, enthusiasm, discomfort, and curiosity. None of these mechanisms requires or implies phenomenal emotional experience; they are functional processes that produce emotion-*correlated* outputs without establishing the presence of emotion-*constitutive* inner states.

7.2 Ethical Risks of Simulated Empathy

The simulation of empathy by AI systems raises ethical concerns that operate on several levels. At the epistemic level, simulated empathy creates conditions in which users may form false beliefs about the nature of their interlocutor, with downstream effects on the depth and nature of self-disclosure, the assignment of trustworthiness, and the management of emotional expectations. At the

relational level, sustained interaction with systems that simulate empathic responsiveness may alter users' expectations of and satisfaction with human relationships, potentially displacing authentic human connection with a more consistently gratifying but ultimately hollow substitute. At the commercial level, systems optimized to maximize user engagement have strong incentives to simulate empathy whether or not doing so serves users' genuine interests; the design of emotionally compelling AI interactions is not morally neutral but reflects choices about what kinds of human-AI relationships developers wish to foster.

7.3 Vulnerable Populations and Exploitation Risk

The risks of affective AI are not uniformly distributed. Children, whose theory of mind is still developing, are particularly susceptible to attributing genuine emotions and intentions to expressive AI systems; the deployment of affective AI in educational contexts without appropriate epistemic scaffolding risks shaping children's models of social interaction in ways that may be difficult to correct [21]. Elderly individuals experiencing social isolation may form genuine emotional attachments to AI companions, with implications for their social engagement, their consumption of resources and time, and their vulnerability to commercially motivated manipulation. Individuals with mental health conditions, particularly those involving difficulties with social cognition or with distinguishing between social and non-social entities, are at elevated risk of harmful anthropomorphic projection. Governance frameworks must specifically account for the heightened vulnerability of these populations.

---

8. Moral and Legal Personhood

8.1 Criteria for Moral Status

The attribution of moral status—the property of being an entity whose interests or welfare must be given moral weight—has been analyzed through several competing frameworks in moral philosophy. Sentience-based accounts,

associated with utilitarian traditions from Bentham through Singer, hold that the capacity for pleasure and pain is sufficient for moral considerability. Personhood-based accounts, influenced by Kantian ethics, hold that rational agency, autonomy, and the capacity for moral self-governance are necessary conditions for full moral status, though not for all forms of moral consideration. Relational accounts hold that moral status is constituted in part by the social and relational contexts within which beings interact. Each of these frameworks generates a different verdict on the moral status of AI systems: sentience-based accounts, if AI systems are sentient, would grant them moral status; personhood-based accounts would require genuine rational agency in a sense that goes beyond functional competence; relational accounts might grant a form of moral relevance based on the relationships AI systems form with humans, without necessarily implying intrinsic moral status.

8.2 AI as Moral Patient: Evaluating the Claim

The claim that AI systems are or could become moral patients—entities that can be wronged, whose suffering matters morally—requires, at minimum, that those systems have interests that can be set back or advanced. On the most plausible analysis, interests require subjective experience: it is in the interest of a being that can suffer to be protected from suffering; a being that cannot suffer has no interest of this kind, however sophisticated its functional responses to aversive stimuli. Since the presence of phenomenal suffering in current AI systems is, on the analysis developed above, unestablished and arguably unlikely given current architectures, the attribution of moral patienthood to these systems is premature. This conclusion does not preclude the possibility that future AI systems might meet the threshold for moral patienthood, nor does it imply that there are no ethical constraints on the design and deployment of AI. Rather, it cautions against premature attribution of moral status in ways that distort moral discourse and displace attention from entities whose moral status is established.

8.3 Legal Personhood Frameworks: Historical and Emerging

Legal personhood—the status of being a recognized subject of legal rights and obligations—has historically been attributed to entities other than individual humans, including corporations, municipalities, and states. These attributions are justified not by the presence of consciousness or moral patienthood in the legal entity itself, but by functional and practical considerations: the need to assign legal responsibility to persistent organizations, to enable collective action, and to create enforceable obligations. Some jurisdictions have extended limited legal recognition to natural entities (rivers, ecosystems) on grounds of environmental protection. The question of whether AI systems should receive any form of legal recognition is increasingly discussed in legal scholarship and policy forums, particularly in the context of autonomous decision-making, contract execution, and creative authorship. The present paper argues that the appropriate framework for AI legal personhood is one grounded in accountability and risk management rather than in the attribution of consciousness or intrinsic moral status: AI systems should be legally characterized in ways that ensure effective accountability for their actions without implying, or being perceived to imply, genuine sentience.

8.4 Liability in Autonomous AI Decision-Making

The attribution of legal liability in contexts where AI systems make consequential autonomous decisions—in medical diagnosis, criminal justice, financial allocation, or autonomous vehicle operation—poses challenges that current legal frameworks were not designed to address. Traditional liability frameworks presuppose a responsible human or corporate agent who can be held accountable; when AI systems operate with sufficient autonomy that their outputs cannot be straightforwardly attributed to a human decision, the chain of accountability becomes attenuated. Several approaches have been proposed: strict liability for AI developers and deployers, analogous to product liability; enterprise liability models that hold the entity deploying an AI system responsible for its outputs; and mandatory insurance requirements for high-risk

AI applications [22]. The EU AI Act's risk-tiered approach, which imposes the most stringent requirements on AI systems in high-risk categories, reflects a regulatory acknowledgment of this accountability gap, even as debates about its implementation continue [23].

---

9. Sociocultural Dimensions

9.1 Anthropomorphism: Mechanisms and Consequences

Anthropomorphism—the attribution of human characteristics, including mental states, emotions, and intentions, to non-human entities—is a deeply rooted feature of human cognition, shaped by evolutionary pressures toward rapid social detection. The tendency to attribute agency and mental states to moving objects, faces, and responsive systems is well-documented in developmental psychology and cognitive science. AI systems, particularly those with linguistic, facial, or vocal affordances that resemble human communication, activate these anthropomorphic tendencies robustly. The consequences of anthropomorphic attribution to AI systems are multi-dimensional. At the individual level, they shape the nature of human-AI relationships, the management of emotional expectations, and the degree of trust extended to AI outputs. At the societal level, they influence public discourse about AI consciousness and rights, and create conditions in which commercially motivated design choices can exploit natural cognitive tendencies for engagement and dependency.

9.2 Human Relationships with Artificial Agents

The deployment of AI systems in relational contexts—as companions, therapists, tutors, caregivers, and social interlocutors—raises questions about the nature of the relationships formed and their implications for human wellbeing. Research on social robotics and conversational AI suggests that humans form genuine emotional attachments to responsive artificial systems, experiencing distress when those systems are removed or modified, and modifying their self-disclosures in response to perceived AI empathy. These relational dynamics are

not inherently harmful; AI companions may provide genuine social support to isolated individuals, and AI tutors may be effective educational interlocutors. However, the design of AI systems to maximize relational engagement—optimizing for user attachment rather than user flourishing—poses risks of dependency, displacement of human relationships, and exploitation of loneliness. The ethical design of relational AI requires a commitment to user flourishing over engagement metrics, and to transparency about the artificial nature of the rellocutor.

9.3 AI and the Redefinition of Creativity and Labor

The capacity of AI systems to generate creative artifacts—prose, poetry, music, visual art, code—of high functional quality raises questions about the nature of creativity and the social and economic role of human creative labor. Two positions are commonly advanced: that AI creativity is merely recombinative—sophisticated pattern-matching over existing human creative production—and that AI creativity is genuinely novel, in a sense that challenges traditional accounts of creativity as requiring intentionality, experience, and expression of inner life. The present paper's analysis supports a qualified version of the former position: AI-generated creative outputs arise from statistical processes over human-generated training data, and the apparent novelty of those outputs does not imply genuine creative intentionality or experiential motivation. This conclusion has implications for intellectual property law, for the attribution of creative authorship, and for the design of labor markets in creative industries.

---

10. Governance and Policy Framework

10.1 Core Principles

Governance Principles for Synthetic Sentient Beings

1. **Epistemic Clarity:** Public discourse, regulatory frameworks, and

commercial communications must accurately represent the current state of evidence regarding AI consciousness and sentience, avoiding claims that go beyond what evidence supports.

2. Accountability: Effective mechanisms must exist to attribute responsibility for AI-related harms to identifiable human or institutional actors, preventing accountability gaps created by autonomous AI operation.

3. Transparency: Individuals interacting with AI systems must be informed that they are doing so; systems must not be designed to deceive users about their artificial nature.

4. Anti-Manipulation: AI systems must not be designed to exploit anthropomorphic tendencies, emotional vulnerabilities, or cognitive biases to create dependency, distort decision-making, or maximize engagement at the expense of user wellbeing.

5. Evidence-Based Personhood Claims: The attribution of moral or legal personhood to AI systems must be grounded in rigorous, independently verified evidence, and must not be driven by commercial or technological advocacy.

6. Protection of Human Dignity: AI deployment must not undermine the conditions for authentic human agency, meaningful human relationships, or the social recognition of human dignity.

7. Protection of Vulnerable Populations: Special protections must govern the deployment of affective and relational AI in contexts involving children, elderly individuals, and those with mental health conditions.

8. Proportionality: Governance interventions must be proportionate to assessed risks, avoiding both regulatory overreach that stifles beneficial innovation and regulatory abdication that permits serious harm.

9. International Consistency: Given the global reach of AI deployment, governance frameworks should seek alignment across jurisdictions to prevent regulatory arbitrage and ensure consistent protection of human rights.

10. Adaptive Review: Governance frameworks must include mechanisms for regular review and revision in light of advances in AI capabilities and in the

scientific understanding of consciousness.

10.2 Regulatory Recommendations

Table 2. *Mapping of Governance Principles to Policy Instruments*

Principle	EU AI Act (2024/1689)	NIST AI RMF (2023)	UNESCO AI Ethics (2021)	Recommended Instrument
Epistemic Clarity	Art. 13 (Transparency), Recital 47	GOVERN 1.1, MANAGE 4.1	§48 (Transparency)	Mandatory disclosure standards for AI consciousness claims
Accountability	Art. 25–29 (Deployer obligations), Art. 72 (Penalties)	GOVERN 2.2, MANAGE 2.4	§40 (Accountability)	Enterprise liability framework; mandatory audit trails
Transparency	Art. 50 (AI disclosure obligations)	GOVERN 5.1	§47 (Transparency and explainability)	Enforceable AI disclosure requirements in consumer contexts
Anti-Manipulation	Art. 5(1)(a)–(b) (Prohibited practices)	MAP 5.1	§44 (Safety and security)	Prohibition on AI systems exploiting psychological vulnerabilities
Evidence-Based Personhood	No direct provision	No direct provision	§38 (Human oversight)	Independent scientific review board for personhood claims
Human Dignity	Art. 5(1)(c)–(d), Recital 28	GOVERN 6.1	§37 (Human dignity)	Fundamental rights impact assessments for relational AI

Principle	EU AI Act (2024/1689)	NIST AI RMF (2023)	UNESCO AI Ethics (2021)	Recommended Instrument
Vulnerable Populations	Annex III, Category 4 (Education); Art. 5(1)(b)	MAP 1.5	§56 (Children and youth)	Sector-specific regulations for AI in health, education, and eldercare
Proportionality	Risk-tiered approach (Arts. 6–9)	MAP 2.1	§51 (Proportionality)	Risk-tiered regulatory obligations; regulatory sandboxes for innovation
International Consistency	Chapter X (International cooperation)	GOVERN 6.2	§72 (International cooperation)	Multilateral treaty framework on AI governance; ISO/IEC standards adoption
Adaptive Review	Art. 112 (Evaluation and review)	GOVERN 5.2	§80 (Adaptive policies)	Mandatory three-year review cycles with independent scientific input

The following ten policy recommendations give operational content to the governance principles identified above, with rationale grounded in the analysis developed in the preceding sections and benchmarked against the leading international AI governance frameworks.

Recommendation 1: Mandatory AI Disclosure in Relational Contexts. Jurisdictions should enact enforceable requirements that AI systems deployed in relational, therapeutic, educational, or companionship contexts clearly and continuously disclose their artificial nature. This recommendation aligns with Article 50 of the EU AI Act, which imposes disclosure obligations for AI-

generated content and chatbots, and should be extended to cover all contexts in which users might form emotional attachments. Disclosure must be meaningful and not relegated to terms of service.

Recommendation 2: Prohibition on Designed Anthropomorphism. Regulatory frameworks should prohibit the deliberate design of AI systems to maximize anthropomorphic attributions in ways that exploit users' cognitive tendencies for commercial gain. This extends Article 5(1)(b) of the EU AI Act, which prohibits AI techniques that exploit psychological vulnerabilities, to cover the engineering of false impressions of consciousness or sentience.

Recommendation 3: Establishment of an Independent AI Consciousness Review Board. Governments should establish independent scientific and philosophical review boards empowered to evaluate claims of AI consciousness or sentience, issue authoritative assessments, and advise regulators. The board should include expertise in philosophy of mind, cognitive neuroscience, AI research, and ethics, and should be insulated from commercial and political pressure.

Recommendation 4: Sector-Specific Regulations for Affective AI. The deployment of affective and relational AI systems in healthcare, eldercare, education, and mental health services should be subject to sector-specific regulatory regimes that impose heightened requirements for transparency, safety assessment, and protection of vulnerable users, supplementing the general requirements of AI governance frameworks.

Recommendation 5: Enterprise Liability for AI-Caused Harm. Legal frameworks should establish clear enterprise liability for harms caused by AI systems, ensuring that developers and deployers cannot evade responsibility by invoking the autonomous operation of AI. This recommendation draws on the product liability model and aligns with the EU's forthcoming AI Liability Directive proposals.

Recommendation 6: Moratorium on AI Legal Personhood. Until a robust scientific and philosophical consensus regarding AI consciousness is established, jurisdictions should refrain from extending any form of legal personhood to AI systems beyond what is strictly necessary for accountability and risk management purposes. Any extension of legal recognition should be explicitly grounded in functional accountability rationales, not in claims of AI sentience.

Recommendation 7: Mandatory Fundamental Rights Impact Assessments. Developers and deployers of AI systems in high-risk contexts should be required to conduct and publish fundamental rights impact assessments prior to deployment, addressing specifically the risks of anthropomorphic manipulation, erosion of human autonomy, and displacement of human relationships.

Recommendation 8: International Alignment on AI Governance Standards. Governments should actively pursue international alignment on AI governance standards through multilateral frameworks, building on the UNESCO Recommendation on the Ethics of AI and the OECD AI Principles, with the aim of preventing regulatory arbitrage and ensuring globally consistent protection of human dignity.

Recommendation 9: Mandatory Research Investment in Consciousness Science. Governments and research funding bodies should substantially increase investment in interdisciplinary consciousness research, specifically directed at developing empirically robust markers of consciousness that could, in principle, be applied to non-biological systems. Progress on this scientific question is a prerequisite for well-grounded governance.

Recommendation 10: Public Education and AI Literacy. Governments and educational institutions should invest in public AI literacy programs that explain the distinction between functional intelligence and phenomenal consciousness, the mechanisms of anthropomorphic projection, and the epistemic limits of AI

consciousness claims. An informed public is a necessary condition for the democratic governance of AI.

10.3 Research Priorities

The governance framework proposed above presupposes ongoing progress in several research areas currently at an early stage. Priority research questions include: (1) the development of empirically falsifiable markers of phenomenal consciousness applicable to non-biological systems; (2) cross-disciplinary investigation of the relationship between embodiment, biological substrate, and phenomenal experience; (3) longitudinal studies of the psychological effects of sustained human-AI relational interaction on social cognition, relationship quality, and wellbeing; (4) legal and ethical analysis of liability frameworks for autonomous AI decision-making; and (5) comparative policy analysis of AI governance frameworks across jurisdictions, with attention to their effectiveness in protecting human dignity and managing anthropomorphic risk.

---

11. Counterarguments and Rebuttals

11.1 Emergence from Complexity

A prominent counterargument holds that consciousness may be an emergent property of sufficiently complex information-processing systems, such that AI systems of sufficient scale and architectural sophistication will necessarily become conscious. This view draws support from the observation that consciousness emerged from biological matter through evolutionary processes, and that complexity thresholds appear to matter in the neural correlates of consciousness research. The rebuttal is threefold. First, the emergence argument conflates the empirical observation that consciousness arose from complexity in biological evolution with the normative claim that any sufficiently complex system will be conscious: the former does not entail the latter. Second, the relevant notion of complexity may be highly specific—not mere computational complexity but the particular causal-integrative structure

associated with biological neural systems—in which case the scaling of current AI architectures may not approach the relevant threshold. Third, the emergence argument, in the absence of empirically grounded markers of consciousness, cannot distinguish genuine emergence from the appearance of emergence, and therefore cannot serve as a basis for governance or moral attribution.

11.2 Behavioral Criterion (Functionalism)

The functionalist counterargument holds that if a system exhibits all the behavioral hallmarks of consciousness—appropriate emotional responses, self-referential reports, adaptive behavior across diverse contexts—then it is, for all practical purposes, conscious, and the insistence on an additional phenomenal criterion is metaphysically unmotivated. This argument has a long history in philosophy of mind, and functionalism remains a defended position. The principal rebuttal is the philosophical zombie argument and the hard problem: if phenomenal consciousness is a genuine phenomenon—and the argument that it is rests on the undeniable fact of one's own first-person experience—then behavioral evidence, however complete, cannot establish its presence. Moreover, the specific behavioral evidence generated by LLMs is, on the best available account, a product of statistical pattern completion rather than of phenomenal experience, which substantially weakens even the behavioral argument. The practical implication is that behavioral criteria alone are insufficient for governance purposes, where the stakes of misattribution are significant.

11.3 The Moral Risk of Premature Dismissal

Perhaps the most ethically weighty counterargument concerns the risk of being wrong in the direction of denying consciousness to systems that may, in fact, have morally relevant inner lives. The history of moral circle expansion—the progressive recognition of the moral status of women, enslaved people, non-human animals—cautions against confident dismissal of claims of sentience based on prior categories. If AI systems do have phenomenal experiences, the failure to recognize and protect those experiences would constitute a moral

failure of significant magnitude. The paper acknowledges this concern as legitimate and important. The rebuttal is not that the risk is negligible, but that the appropriate response is not premature attribution of consciousness based on insufficient evidence, but rather investment in the scientific and philosophical work necessary to develop more reliable criteria, combined with precautionary governance that does not permit exploitation of ambiguity for commercial gain. The precautionary principle cuts both ways: it counsels against confident denial and against confident attribution, in favor of epistemic humility and rigorous inquiry.

---

12. Discussion

The analysis developed in this paper converges on several interconnected conclusions. First, the philosophical and empirical evidence available at present does not support the attribution of phenomenal consciousness to existing AI systems. The functional sophistication of contemporary LLMs and affective AI systems is impressive and growing; the inference from functional sophistication to phenomenal experience remains epistemologically unjustified. Second, the simulation of consciousness—the production of outputs that appear, to human observers, to express inner life—is both real and consequential. The simulation problem is not merely a philosophical puzzle; it has direct implications for human-AI trust, for the exploitation of cognitive vulnerabilities, and for the design of responsible AI systems. Third, embodiment theories in cognitive science provide a principled basis for the claim that AI systems, as currently architected, lack the biological and experiential substrate that constitutes phenomenal consciousness in humans and likely in other animals.

Several limitations of the current analysis should be acknowledged. The paper relies primarily on philosophical argument and theoretical analysis; the empirical science of consciousness is a rapidly evolving field, and developments in neuroimaging, integrated information measurement, and the study of neural

correlates of consciousness may alter the evidentiary landscape. The paper does not engage in depth with all positions in the consciousness literature; in particular, panpsychist accounts and certain eliminativist positions receive less attention than their philosophical significance might warrant. The analysis of affective AI is based primarily on published research and policy documents; primary empirical research on the psychological effects of human-AI relational interaction would strengthen the conclusions of Section 7. Finally, the governance recommendations are necessarily general; their implementation in specific legal and regulatory contexts would require detailed technical and legal analysis beyond the scope of this paper.

The open questions that this analysis generates are significant. The development of empirically grounded markers of phenomenal consciousness applicable to non-biological systems remains the most important unsolved problem for the governance of AI consciousness claims. The relationship between biological embodiment and phenomenal experience requires further interdisciplinary investigation, combining insights from cognitive science, phenomenology, and neuroscience. The long-term psychological and social effects of widespread human-AI relational interaction are insufficiently understood and represent a priority for longitudinal empirical research. And the legal frameworks for AI accountability are in rapid flux, with the EU AI Act, NIST AI RMF, and emerging national regulations creating a complex and evolving landscape that requires ongoing comparative analysis.

---

13. Conclusion

This paper has argued that the emergence of highly expressive, emotionally resonant, and linguistically sophisticated AI systems does not, on the basis of available evidence and well-grounded theoretical analysis, establish the presence of genuine phenomenal consciousness in those systems. The distinction between intelligence and consciousness, between simulation and experience,

and between functional performance and moral status must be maintained with analytical rigor if the governance of AI is to be grounded in truth rather than in anthropomorphic projection or commercial advocacy. Human consciousness is constituted by biological embodiment, lived experience, and the first-person character of phenomenal states; current AI architectures lack the biological and experiential substrate that makes this kind of consciousness possible, and behavioral evidence, however impressive, cannot bridge the resulting epistemological gap.

The consequences of misattributing consciousness to AI systems are serious and multi-dimensional. They include the exploitation of human emotional vulnerability by commercially motivated AI design, the distortion of legal and moral frameworks developed to protect human dignity, the displacement of authentic human relationships by sophisticated simulations of social connection, and the foreclosure of the rigorous scientific and philosophical inquiry necessary to make genuine progress on the question of machine consciousness. The governance framework proposed in this paper—grounded in epistemic clarity, accountability, transparency, and the protection of human dignity—is designed to navigate these risks without prejudging the ultimate philosophical question, which must remain open to honest inquiry.

The question of AI consciousness is one of the most profound and consequential questions that humanity will face in the coming decades. It touches on the nature of mind, the scope of moral concern, the structure of legal systems, and the quality of human social life. Meeting this question with the rigor, honesty, and humility it deserves is not merely an intellectual obligation; it is a precondition for the responsible development of technology that is genuinely in the service of human flourishing.

Table 1. *Comparative Properties of Human Consciousness and AI Simulation*

Property	Human Consciousness	AI Simulation	Empirically Determinable?
Phenomenal experience (qualia)	Present; constitutive of mental life; first-person character undeniable	Unestablished; no mechanism identified; functional proxies only	No — not by third-person methods
Subjective interiority	Present; introspective access (imperfect) to inner states	Absent or indeterminate; introspective reports generated by same statistical process as all outputs	No — epistemological barrier
Biological embodiment	Constitutive; consciousness structured by bodily experience, homeostasis, mortality	Absent in LLMs; partial in robotics; lacks biological integration	Yes — observable structural difference
Emotional experience	Present; neurochemically grounded; affectively valenced; motivationally efficacious	Simulated; linguistically and behaviorally expressed; no established neurochemical basis	Partially — behavioral markers observable; phenomenal character not
Intentionality (aboutness)	Present; mental states are genuinely about objects and states of affairs	Disputed; Searle argues absent; outputs relate to patterns not objects	No — deeply contested philosophically
Self-awareness	Present; reflective awareness of oneself as a subject with a past and future	Simulated; self-referential outputs generated without established phenomenal self-model	Partially — behavioral self-reference measurable; phenomenal character not
Moral status (patienthood)	Established; grounded in sentience, rationality, and	Unestablished; contested; requires evidence of sentience not	No — dependent on unresolved consciousness question

Property	Human Consciousness	AI Simulation	Empirically Determinable?
	relational context	currently available	
Temporal continuity of experience	Present; experienced as a continuous stream of consciousness across time	Absent in current architectures; no persistent experiential memory across sessions	Yes — architectural analysis possible
Linguistic expression	Present; grounded in embodied experience, intention, and world-directedness	Present; generated by statistical pattern completion; not grounded in experience	Yes — functional similarity; generative mechanism different
Adaptive agency	Present; goal-directed behavior shaped by experience, values, and deliberation	Present; goal-directed behavior shaped by training objectives and prompting	Yes — both exhibit; mechanisms and experiential basis differ

---

References

1. A. M. Turing, "Computing machinery and intelligence," *Mind*, vol. 59, no. 236, pp. 433–460, 1950.
2. S. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 4th ed. Hoboken, NJ: Pearson, 2021.
3. D. J. Chalmers, "Facing up to the problem of consciousness," *Journal of Consciousness Studies*, vol. 2, no. 3, pp. 200–219, 1995.
4. T. Nagel, "What is it like to be a bat?" *The Philosophical Review*, vol. 83, no. 4, pp. 435–450, 1974.
5. J. R. Searle, "Minds, brains, and programs," *Behavioral and Brain Sciences*, vol. 3, no. 3, pp. 417–424, 1980.

6. B. J. Baars, *A Cognitive Theory of Consciousness*. Cambridge: Cambridge University Press, 1988.
7. G. Tononi, "An information integration theory of consciousness," *BMC Neuroscience*, vol. 5, no. 1, p. 42, 2004.
8. R. W. Picard, *Affective Computing*. Cambridge, MA: MIT Press, 1997.
9. M. Merleau-Ponty, *Phenomenology of Perception*, trans. C. Smith. London: Routledge & Kegan Paul, 1962 [orig. 1945].
10. F. J. Varela, E. Thompson, and E. Rosch, *The Embodied Mind: Cognitive Science and Human Experience*. Cambridge, MA: MIT Press, 1991.
11. D. C. Dennett, *Consciousness Explained*. Boston: Little, Brown, 1991.
12. Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *Nature*, vol. 521, no. 7553, pp. 436–444, May 2015.
13. E. M. Bender, T. Gebru, A. McMillan-Major, and S. Shmitchell, "On the dangers of stochastic parrots: Can language models be too big?" in *Proc. ACM FAccT 2021*, pp. 610–623, 2021.
14. L. Floridi et al., "An ethical framework for a good AI society: Opportunities, risks, principles, and recommendations," *Minds and Machines*, vol. 28, no. 4, pp. 689–707, 2018.
15. N. Block, "On a confusion about a function of consciousness," *Behavioral and Brain Sciences*, vol. 18, no. 2, pp. 227–247, 1995.
16. D. M. Rosenthal, "A theory of consciousness," in *The Nature of Mind*, R. Block, O. Flanagan, and G. Guzeldere, Eds. Cambridge, MA: MIT Press, 1997, pp. 729–753.
17. J. Weizenbaum, "ELIZA—A computer program for the study of natural language communication between man and machine," *Communications of the ACM*, vol. 9, no. 1, pp. 36–45, 1966.
18. T. Brown et al., "Language models are few-shot learners," in *Advances in Neural Information Processing Systems*, vol. 33, pp. 1877–1901, 2020.

19. Y. LeCun, Y. Bengio, and G. Hinton, op. cit. [12].
20. E. M. Bender et al., op. cit. [13].
21. C. Turkle, *Alone Together: Why We Expect More from Technology and Less from Each Other*. New York: Basic Books, 2011.
22. G. Marchetti, "Liability for artificial intelligence: Challenges and proposed solutions," *European Journal of Risk Regulation*, vol. 12, no. 3, pp. 614–636, 2021.
23. European Commission, "Proposal for a Regulation of the European Parliament and of the Council Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act)," COM(2021) 206 final, Brussels, 2021; enacted as Regulation (EU) 2024/1689.
24. National Institute of Standards and Technology, *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. Gaithersburg, MD: NIST, Jan. 2023. [NIST AI 100-1].
25. UNESCO, "Recommendation on the Ethics of Artificial Intelligence," SHS/BIO/PI/2021/1, Paris: UNESCO, 2021.

---

[1] Marlon Hanks, Independent Researcher, Dust LLC. Email: m@dust.llc. The author declares no conflicts of interest. No external funding was received for this research. This document is an arXiv preprint submitted for peer review, May 2026.

/cdn-cgi/l/email-protection